

LG-CISH: Result Tables and Matching Figures

Empirical Evaluation of Coverless Image Steganography (codebook $N=40$)

Ayushman Bhattacharya Soumyashree Biswas Piyali Sarkar Harsh Raj

All numbers are produced by the reproducible evaluation suite (`python evaluation/generate_all.py`) on a codebook of $N=40$ visually distinct 512×512 images (UCID / Kodak / USC-SIPI), $\log_2 40 \approx 5.322$ bits/image. Each table is shown with the figure that visualises it.

1 Experimental Setup

Table 1: Experimental configuration.

Parameter	Value
Dataset	UCID + Kodak + USC-SIPI (mixed), 512×512 PNG
Codebook images N	40
Capacity	5.322 bits/image (base-40)
CLIP model	ViT-B/32, $d=512$
LLM (plausibility / aliases)	Gemini 2.5 Flash-Lite (vision) + FLUX (gen)
Coding modes	base- N (5.322 b/img) ; permutation (no-repeat)
Separation threshold τ	0.85
Pairwise sim. (min/mean/max)	0.305 / 0.513 / 0.827
Decoding margin ($1 - \max$)	0.173
Encryption / integrity / compression	AES-256-CBC / CRC-32 / zlib
Hardware	RTX 3050 6 GB, 16-core CPU




Figure 1: Pairwise CLIP cosine similarity of the $N=40$ codebook. All off-diagonal values stay below $\tau=0.85$ (max 0.827), giving a positive separation margin for nearest-neighbour recovery (matches Table 1).

2 Qualitative Results




Figure 2: Figure Set 1 — secret message → generated image sequence (short / medium / long). The identity and order of unmodified images encode the data.



Figure 3: Figure Set 2 — received sequence → CLIP nearest-neighbour decode; original and reconstructed messages match exactly.




Figure 4: Figure Set 3 — failure analysis. Accuracy collapses only under extreme degradation; a codebook mismatch is rejected by CRC-32.

3 Quantitative Evaluation

Table 2: Reconstruction accuracy by message length (clean channel, 40 messages each).

Message length	Msgs	Accuracy (%)	BER	Hash match (%)	Avg images
Short (≤ 50 chars)	40	100.00	0	100.00	57.5
Medium (50–200)	40	100.00	0	100.00	153.2
Long (200–1000)	40	100.00	0	100.00	399.4

Table 3: Payload capacity vs. distortion. LG-CISH trades raw capacity for zero distortion.

Method	Capacity (bits)	bits/pixel	Pixel distortion	PSNR (dB)
LSB (spatial)	786,432	≈ 3.0	Yes (high)	51.1
DCT-LSB	4,096	≈ 0.016	Yes (moderate)	38–42
DWT-DCT [cited]	$\sim 4,096$	≈ 0.015	Yes (low)	40–44
Proposed LG-CISH	5.322/image	5.322	None (coverless)	∞

Table 4: Coding modes. Permutation coding trades a little capacity for a no-repeat, more plausible cover.

Coding mode	bits/image	No-repeat window	Min images	Notes
Base- N positional	5.322	none (repeats allowed)	1	max capacity
Permutation (block = 8)	5.187	8 images	8	no repeats / block
Permutation (block = 16)	5.008	16 images	16	no repeats / block
Permutation (block = 40= N)	3.979	40 images	40	full permutation




Figure 5: Capacity per image for base- N vs. permutation coding at several block sizes (matches Table 4). Smaller permutation blocks recover most of the base- N ceiling while guaranteeing no in-block image repeats.

Table 5: Computational time (mean over repeated runs, RTX 3050).

Stage	Time (ms)
Full encode (200-char message)	0.04
CLIP embedding (per image)	15.36
Index recovery (189 images)	1873.31
Full decode (189 images)	1816.14
Decode per image (amortised)	9.61




Figure 6: Per-stage timing on log scale (matches Table 5). Encoding is sub-millisecond; cost is dominated by CLIP embedding at decode.




Figure 7: Reconstruction accuracy vs. payload length: 100% across all sizes while the image count grows linearly (matches Table 2).




Figure 8: Capacity (bits/pixel) vs. distortion (PSNR). LG-CISH sits at infinite PSNR — no pixel is modified (matches Table 3).

4 Robustness Analysis

Table 6: Robustness to all channel attacks (40 messages/attack).

Attack	Accuracy (%)	BER	CLIP margin	Top-1 sim
No attack (baseline)	100.00	0	0.302	1.000
JPEG 90%	100.00	0	0.299	0.985
JPEG 70%	100.00	0	0.294	0.976
JPEG 50%	100.00	0	0.289	0.961
JPEG 30%	100.00	0	0.280	0.942
JPEG 20%	100.00	0	0.270	0.930
Gaussian $\sigma=5$	100.00	0	0.293	0.986
Gaussian $\sigma=10$	100.00	0	0.285	0.974
Gaussian $\sigma=20$	100.00	0	0.272	0.957
Gaussian $\sigma=30$	100.00	0	0.260	0.942
Salt & Pepper 0.01	100.00	0	0.277	0.960
Salt & Pepper 0.05	100.00	0	0.246	0.922
Salt & Pepper 0.10	100.00	0	0.214	0.882
Resize 50%	100.00	0	0.282	0.982
Resize 25%	100.00	0	0.246	0.939
Resize 10%	12.50	1.08×10^{-2}	0.125	0.790
Crop 95%	100.00	0	0.280	0.979
Crop 90%	100.00	0	0.268	0.969
Crop 85%	100.00	0	0.261	0.961
Crop 60%	100.00	0	0.196	0.910
PNG→WebP 80	100.00	0	0.295	0.976




Figure 9: Mean decoding margin $\gamma = s_{(1)} - s_{(2)}$ per attack. All but the most extreme downscale leave $\gamma > 0$ (matches Table 6).




Figure 10: Decoding accuracy vs. JPEG quality: LG-CISH stays at 100% where pixel-domain LSB/DCT collapse.

5 Security and Steganalysis Resistance

Table 7: Chi-square steganalysis detection ($\approx 50\%$ = chance). $N=20$ patches.

Method	Detection acc. (%)	TPR (%)	TNR (%)
LSB	97.5	100.0	95.0
DCT-LSB	47.5	0.0	95.0
Proposed LG-CISH	47.5	0.0	95.0

Table 8: Keyspace analysis ($N=40$).

Component	Size
Codebook orderings	$40! \approx 2^{159.2}$
AES-256 key	2^{256}
Combined keyspace	$\approx 2^{415.2}$

Table 9: Cover plausibility (Gemini judge, 0–1). Plausibility is driven by codebook theme, not the coding mode.

Configuration	Plausibility (0–1)	Note
Diverse benchmark (UCID/Kodak/USC)	0.12 ± 0.04	mixed subjects look random
Themed codebook (coherent context)	0.88 ± 0.07	looks like a personal album
Diverse + permutation coding	0.18 ± 0.10	codec barely moves the score




Figure 11: Measured steganalysis detection rates (lower is stealthier). LG-CISH sits at the chance line, far below LSB (matches Table 7).

6 Ablation Study

Table 10: Ablation. Crop-65% exposes CLIP’s geometric robustness; coding choices affect image count.

Variant	Clean (%)	JPEG50 (%)	Crop65 (%)	Avg images	Integrity
Full LG-CISH (proposed)	100	100	100	134.4	CRC-32
Without CLIP (pHash)	100	100	0	134.4	CRC-32
Without compression	100	100	100	182.2	CRC-32
Without CRC (no integrity)	100	100	100	134.4	None
Fixed-chunk (5-bit)	100	100	100	143.0	CRC-32




Figure 12: Left: base- N + zlib needs the fewest images. Right: under Crop-65% CLIP keeps 100% while pHash drops to 0% (matches Table 10).

7 Comparative Analysis

Table 11: Comparison with classical and coverless baselines. Bracketed values are cited.

Method	Capacity	JPEG50 (%)	Detection (%)	Time	PSNR (dB)
LSB	786,432 bits	51.4	98	0.01–1	51.1
DCT-LSB	~4096 bits	42.1	[70–90]	1–5	38–42
DWT-DCT [cited]	~4096 bits	[85]	[60–80]	~50	40–44
Coverless [cited]	low	[95]	[~50]	high	∞
Proposed LG-CISH	5.322 bits/img	100.0	48	fast	∞

8 Statistical Validation

Table 12: Mean $\pm$ std and 95% confidence intervals over $N=50$ trials (clean channel).

Metric	Mean $\pm$ Std	95% CI
Reconstruction accuracy (%)	100.00 $\pm$ 0.00	[100.00, 100.00]
Bit error rate	0 $\pm$ 0	[0, 0]
CLIP margin	0.3020 $\pm$ 0.0059	[0.3003, 0.3037]
Images per message	181.2 $\pm$ 44.5	[168.5, 194.0]

Table 13: Significance of CLIP’s geometric robustness over pHash: paired Wilcoxon signed-rank across a 7-level crop sweep. Proposed-vs-LSB @ JPEG-50 shown descriptively.

Group / statistic	Accuracy (%)	Std / note
Proposed (CLIP) — crop-sweep mean	100.0	0.0
pHash ablation — crop-sweep mean	0.0	0.0
Wilcoxon W (paired, 7 levels)	28.0	—
p -value (one-sided, CLIP > pHash)	7.81×10^{-3} (< 0.05)	
Proposed vs. LSB @ JPEG-50	100 vs 50	descriptive